

CSE 4/546: Reinforcement Learning

Spring 2022

Instructor: Alina Vereshchaka

Assignment 2 - Deep Q-Networks

Checkpoint: March 27, Sun, 11:59pm

Due Date: April 10, Sun, 11:59pm

1 Assignment Overview

The goal of the assignment is to work with value function approximation algorithms, to explore OpenAI Gym environments and getting experience working with a project-management tool. In the first part of the project we will explore OpenAI gym library. In the second part we will implement Deep Q-learning (DQN) following DeepMind's paper that explains how reinforcement learning algorithms can learn to play Atari from raw pixels. In the third part of the project we will implement an improvement to the DQN algorithm. The purpose of this assignment is to understand the benefits of approximation methods, the role of deep neural networks as well as some of the techniques used in practice to stabilize training and to achieve better performance. We will train our networks on the grid-world and OpenAI gym environments.

For this assignment, libraries with in-built RL methods cannot be used. Submissions with used in-built RL methods (e.g. stable-baselines, keras-RL, TF agents, etc) will not be evaluated.

Part I [10 points] - Exploring OpenAI Gym Environments

Install OpenAI gym library and explore the environments and main functions. and one other environment that you use for the assignment (e.g. possible actions/states, goal, rewards, etc).

In your report for Part I:

- Explore 'CartPole-v1' and provide the main details about the environment (e.g. possible actions/states, goal, rewards, etc).
- Choose one more environment that you use for the assignment and provide the main details about it

Part I submission includes

- Report (as a pdf file)
- Jupyter Notebook (.ipynb) – one file for all explored environments with saved outputs

Part II [40 points] - Implementing DQN & Solving grid-world environment

2.1 Implementing DQN

Implement DQN from scratch following DeepMind's paper ([Nature paper](#) or [initial version](#)). You may use any framework (Keras/Tensorflow/Pytorch).

2.2 Solving grid-world environment

Apply DQN to solve the environment you or your teammate have defined in Assignment 1. You can make slight modifications or improvements to the original environment, if required.

In your report for Part II:

1. Discuss the benefits of:
 - Using experience replay in DQN and how its size can influence the results
 - Introducing the target network
 - Representing the Q function as $\hat{q}(s, \mathbf{w})$
2. Briefly describe the grid-world environment that you used (e.g. possible actions, states, agent, goal, rewards, etc). You can reuse related parts from your Assignment 1 report.
3. Show and discuss your results after applying your DQN implementation on the environment. Plots should include epsilon decay and the reward per episode.
4. Provide the evaluation results. Run your environment for at least 5 episodes, where the agent chooses only greedy actions from the learnt policy. Plot should include the total reward per episode.

Part II submission includes

- Report (as a pdf file) – combined with your report from Part I
- Jupyter Notebook (.ipynb) with all saved outputs. It can be combined with Part I .ipynb

Part III [50 points] - Improving DQN & Solving OpenAI Gym Environments

3.1 Improving vanilla version of DQN

There have been many improvements developed for the vanilla algorithm. In this part we will implement one of improved algorithms that is based on DQN. Modify your DQN code from Part 2.1 to one of the improved versions and apply it to one grid-world environment from Part 2.2 and compare the results.

Possible algorithms to implement include:

- Double DQN
- Dueling DQN
- Prioritized Experience Replay (PER)

3.2 Applying DQN-based algorithms to solve TWO complex environments

Test your DQN algorithm implemented in Part 1 and the improved version on 'CartPole-v1' and a complex environment. You may use your custom made multi-agent environment or any other complex environment that you will use for your Final Project (this has to be approved by the course instructors). Compare the results. Provide reward dynamics (average reward in t -steps). The environment with multiple versions is considered one environment.

'CartPole-v1' CartPole is one of the classical OpenAI gym environments that often used as a benchmark problem to check the algorithm performance. 'CartPole-v1' environment is considered to be solved if it is getting an average reward of more than 470 points over 100 consecutive episodes during evaluation.

Suggested second environment to work with:

- OpenAI LunarLander
- OpenAI Atari Breakout
- OpenAI MountainCar

In your report for Part III:

1. Discuss the algorithm you implemented.
2. What is the main improvement over the vanilla DQN?
3. Show and discuss your results after applying your the two algorithms implementation on the environment. Plots should include epsilon decay and the reward per episode.
4. Provide the evaluation results. Run your environment for at least 5 episodes, where the agent chooses only greedy actions from the learnt policy. Plot should include the total reward per episode.
5. Compare the performance of both algorithms (DQN & Improved version of DQN) on the same environments (e.g. show one graph with two reward dynamics) and provide your interpretation of the results. Overall three rewards dynamics plots with results from two algorithms applied on:
 - Grid-world environment
 - 'CartPole-v1'
 - OpenAI Gym or Multiagent environment
6. Provide your interpretation of the results. E.g. how the same algorithm behaves on different environments, or how various algorithms behave on the same environment.

Part III submission includes

- Report (as a pdf file) – combined with your report from Part I and Part II
- Jupyter Notebook (.ipynb) with all saved outputs. It can be combined with Part I and Part II .ipynb

Extra Points [max +12 points]

- **Using PyTorch [2 points]**

PyTorch is becoming a standard framework to use for research related to AI (e.g [OpenAI](#)), as well becoming more popular in industry. It also mostly provides a better performance and faster convergence on RL-related tasks. We want to motivate you to develop skills working with this framework. If you choose to complete your assignment using PyTorch, you will receive additional points.

- **Solving Image-based Environment [7 points]**

Use one of the environments with image representation of the state that requires to utilize CCN (Convolution Neural Network) for the state preprocessing (e.g. OpenAI Breakout).

- **Project Management [3 points]**

Project management includes the application of specialized skills, knowledge, processes and procedures to successfully take an assignment from start to finish. These tools can be used for collaboration, status reporting, time management and many other cases. There is an extensive range of tools available with free tiers available for small team usage. You can choose any of the following: [Trello](#), [JIRA](#), [Basecamp](#).

Requirements:

- Create a board with at least these columns: To do, In Progress, Done
- Divide the project to at least 20 steps/milestones (E.g. Implement DQN, Prepare grid-world env, Generate rewards dynamic graph, Report for Part 1 Report for Part 2: Q1/Q2/Q2, Submit Checkpoint, Submit Final, etc).
- Add course instructors as observers or as a team member to your board [devikabi@buffalo.edu and nitinvis@buffalo.edu]
- The board should be created at least one week before the checkpoint submission and every week there should be some activities

4 Deliverables

Submit your work using UBLearn's group in both cases if you work individually or in a team of two. There are two parts in your submission:

4.1 Report

The report should be delivered as a separate pdf file, and it is recommended for you to use the NIPS template to structure your report. You may include comments in the Jupyter Notebook, however you will need to duplicate the results in the separate pdf file.

For the final submission, combine the reports for all parts into one file UBIT

TEAMMATE1_TEAMMATE2_assignment2_report.pdf

(e.g. *avereshc_nitinvis_assignment2_report.pdf*)

4.2 Code

Python is the only code accepted for this project. You can submit the code in Jupyter Notebook with the saved results. You can submit multiple files, but they all need to have a clear name. After executing command Jupyter Notebook, it should generate all the results and plots you used in your report and should be able to be printed out in a clear manner. Additionally you can submit the trained parameters, so that the grader can fully replicate your results.

For the final submission we should expect to have Jupyter Notebooks named as

TEAMMATE1_TEAMMATE2_assignment2.ipynb

(e.g. *avereshc_nitinvis_assignment2.ipynb* or *TEAMMATE1_TEAMMATE2_assignment2_partI.ipynb*)

(e.g. *avereshc_nitinvis_assignment2_partI.ipynb*)

5 References

- [NIPS Styles \(docx, tex\)](#)
- [Overleaf](#) (LaTeX based online document generator) - a free tool for creating professional reports
- [GYM environments](#)

- Lecture slides
- [Human-level control through deep reinforcement learning](#)
- [Prioritized Experience Replay](#)
- [Deep Reinforcement Learning with Double Q-learning](#)

6 ASSIGNMENT STEPS

1. Register your team (Due date: March 18)

- You may work individually or in a team of up to 2 people. The evaluation will be the same for a team of any size.
- Register your team at UBLearns (UBlearns > Tools > Groups). In case you joined the wrong group, make a private post on piazza.

2. Submit checkpoint (Due date: March 27)

- Complete Part I and Part II
- Add your pdf and ipynb to zip file with UBIT *TEAMMATE1_TEAMMATE2_assignment2_checkpoint.zip* (e.g. *avereshc_nitinvis_assignment2_checkpoint.zip*)
- Submit at UBLearns > Assignments
- Checkpoint will be evaluated after the final submission

3. Submit final results (Due date: April 10)

- Fully complete all parts of the assignment
- Add your combined pdf and ipynb for Part 1, Part 2 and Part 3 to a zip file *TEAMMATE1_TEAMMATE2_assignment2_final.zip* (e.g. *avereshc_nitinvis_assignment2_final.zip*)
- Submit at UBLearns > Assignments
- The code of your implementations should be written in Python. You can submit multiple files, but they all need to be labeled clearly.
- Your Jupyter notebook should be saved with the results
- Include all the references that have been used to complete the assignment
- In order to certify that you followed Academic Integrity policy while completing the assignment, please include the following statement as a comment block at the beginning of your code. In case you submit multiple files, add this statement at the beginning of each code file submitted:
 "I/we certify that the code and data in this assignment were generated independently, using only the tools and resources defined in the course and that I/we did not receive any external help, coaching or contributions during the production of this work."
- If you are working in a team of two, we expect equal contribution for the assignment. Each team member is expected to make a code-related contribution. Provide a contribution summary by each team member in a form of a table below. If the contribution is highly skewed, then the scores of the team members may be scaled w.r.t the contribution.

Team Member	Assignment Part	Contribution (%)

7 Academic Integrity

This assignment can be completed individually or in a team of two students. The standing policy of the Department is that all students involved in any academic integrity violation (e.g. plagiarism in any way, shape, or form) will receive an F grade for the course. The catalog describes plagiarism as “Copying or receiving material from any source and submitting that material as one’s own, without acknowledging and citing the particular debts to the source, or in any other manner representing the work of another as one’s own.”. Updating the hyperparameters or modifying the existing code is not part of the assignment’s requirements and will result in a zero. Please refer to the [UB Academic Integrity Policy](#).

IMPORTANT NOTE

In order to certify that you followed Academic Integrity policy while completing the assignment, please include the following statement as a comment block at the beginning of your code. In case you submit multiple files, add this statement at the beginning of each code file submitted:

"I certify that the code and data in this assignment were generated independently, using only the tools and resources defined in the course and that I did not receive any external help, coaching or contributions during the production of this work."

Submissions without the academic integrity statement will not be evaluated and will receive a 0.

8 Important Information

This assignment can be completed in groups of two or individually. Teams can not be the same for A2 & A3 assignments.

9 Late Days Policy

You can use up to 5 late days throughout the course toward any assignments’ checkpoint or final submission. You do not have to inform the instructor, as the late submission will be tracked in UBlerns. If you work in teams the late days used will be subtracted from both partners. E.g. you have 4 late days and your partner has 3 days left. If you submit one day after the due date, you will have 3 days and your partner will have 2 days left.

10 Important Dates

March 18, Fri 11:59pm - Register your team at UBlerns > Tools > Teams

March 27, Sun 11:59pm - Checkpoint is Due

April 10, Sun, 11:59pm - Assignment 2 is Due